

1 Abstract

[COMING SOON]

2 Introduction

To make use of the exciting possibilities that the Semantic Web offers, Artificial Intelligence techniques, including Evolutionary methods, have been used in many of its related tasks. Examples of these include query answering[1] and reasoning[?]. However, part of the great potential of the Semantic Web lies in the ability of software agents to use its semantic structure to collect relevant or related data. This paper explores the use of Evolutionary Programming to train such agents. The agents will be given a starting set of triples, and attempt to effectively query a dataset for related information. Entailments gained from the combined results is used as a scoring mechanism for the relevancy of results. First, the method and structure of the agents is discussed. Then, the software and hardware environment used in testing is specified. Finally, the results gained from testing are interpreted, and any conclusions that can be made regarding this method are given.

3 Method

Here, the approach used by the program to evolve the agents query construction is outlined. All agents are given a starting set of triples, and can load a limited amount of additional triples from the dataset each round.

3.1 Algorithm

The program's main structure is that of an Evolutionary Programming algorithm[3]. For mutation, bit flipping is used instead of Gaussian mutation, since the parameter values have no ordering. Initially, the population is filled with randomly generated agents. Each round, all the agents take a number of steps, and afterwards the round the program reads their scores. Agents keep their own score based on the amount entailments they make. Elitist selection is done by eliminating the ten percent worst achievers, and replacing them with the offspring of the ten percent best achievers.

3.2 Agents

Each agent consists of a personal memory, a genotype of two variably sized parameter lists, and a query list. Agents construct their query list by mapping terms from the personal memory into queries according to genotype. In addition to the personal memory, each agent has access to five common ontological terms for use in the queries. The size of the query list generated can differ in accordance with the parameters, but all agents have a shared maximum of results they can load into memory per round. To score, the agents conduct entailment over their memory and count the number of entailments gained. RDFS entailment patterns are used as the basis[7].

4 Environment

4.1 Dataset

The dataset used as a search space for the agents was the union of the DBpedia persondata file[9] together with the DBpedia ontology[10] to allow for more complex entailments. The DBpedia statistics [source] report 831,558 instances of type Person in its dataset.

At each iteration of searching, an agent can load a maximum of a fifty triples into personal memory in addition to the fifty triples of the initial memory already present. Thus, at each step, the search space of all solutions of a step is:

$$S = \{x : 50 \leq |x| \leq 100, I \subseteq x\} \quad (1)$$

Where G is the complete dataset and I is the initial memory.

4.2 Software

The Sesame framework was used both as a datastore and a querying engine. The entire program was written in Java. Jena was used for its entailment engine.

4.3 Hardware

For the experiment, the compute service of the DAS-4 VU cluster was used. The programs threads were run on twenty-four 2.4 Ghz Intel E5620 processors supported by 130GB of memory.

5 Results

The experiment was run for a thousand rounds over a population of 300 agents. Results were gathered on the changes in genotype and behaviour of the agents.

For the performance of agents, first a baseline was measured as the entailments made over a hundred randomly selected triples from the dataset. In a thousand sets of randomly selected triples, the amount of entailments made were almost exclusively a hundred. These are all single argument Type:Resource entailments. In short, practically no interesting entailments could be made. The top agents performed considerably better, with the best agent of each round averaging 152.4 entailments, and the top ten agents of each round averaging 143.24 entailments. Ironically, the total population actually averaged a lower score than the randomized baseline, at 97.83 entailments. This is most likely due to agents conducting queries giving back no results, and thereby actually ending up with less than the maximum hundred triples.

In the behaviour of the agents, two clear trends were identified. The amount of queries conducted per round rose sharply during the experiment, suggesting a strategy of loading heterogeneous results from varying queries was favourable. Furthermore, in the top performing agents. Furthermore, an increase in homogeneity at each round in the term parameters was noticed, suggesting the success of specific memory-to-query mappings.




Figure 1: Query behaviour per round

Surprisingly, these evolutionary changes in behaviour did not lead to a significant change in performance. The amount of entailments made by both the entire population and the top performing agents was assessed. In no

category was any noteworthy improvement gained during the rounds of the experiment.




Figure 2: Entailment performance per round

6 Conclusion

[COMING SOON]

6.1 Related work

Recently, it has been argued that many Semantic Web related tasks can be more successfully approached as optimization problems suitable for Evolutionary and Swarm computing. Specifically, the scalability of Evolutionary solutions has been noted to be required to deal with the massive and opaque nature of Semantic Web data.[2] Evolutionary Computing has already been used to optimize Query Answering, and again advantages in memory usage and scalability have been noted. [1] For the task of appropriate Query Construction, its complexity has been noted and some solutions have been put forward for it regarding keyword search, although these do not utilize Evolutionary methods.[4][5]

References

- [1] Guret, Christophe, et al. *An evolutionary perspective on approximate rdf query answering*. Scalable Uncertainty Management. Springer Berlin Heidelberg, 2008. 215-228.

- [2] Guret, Christophe, et al. "Evolutionary and swarm computing for the semantic web." *Computational Intelligence Magazine*, IEEE 7.2 (2012): 16-31.
- [3] Eiben, Agoston E., and Marc Schoenauer. "Evolutionary computing." *Information Processing Letters* 82.1 (2002): 1-6.
- [4] Zhou, Qi, et al. *SPARK: adapting keyword query to semantic search*. Springer Berlin Heidelberg, 2007.
- [5] Zenz, Gideon, et al. "From keywords to semantic queriesIncremental query construction on the Semantic Web." *Web Semantics: Science, Services and Agents on the World Wide Web* 7.3 (2009): 166-176.
- [6] Antoniou, Grigoris and Van Harmelen, Frank, *A semantic web primer*. MIT press 2004.
- [7] W3C (2014), *RDF Schema 1.1*. Available at: <http://www.w3.org/TR/rdf-schema/> (Accessed: 16 May 2014).
- [8] W3C (2014), *RDF 1.1 Semantics*. Available at: <http://www.w3.org/TR/2014/REC-rdf11-mt-20140225/> (Accessed: 17 May 2014).
- [9] *Persondata*. Available at: http://downloads.dbpedia.org/3.9/en/persondata_en.nt.bz2 (Accessed: 11 May 2014).
- [10] *DBpedia ontology*. Available at: http://downloads.dbpedia.org/3.9/dbpedia_3.9.owl.bz2 (Accessed: 11 May 2014).